

DiagnosisAI - Technical Architecture

Version: 2.0

Last Updated: December 28, 2025

Author: Bharath Nagesh

Table of Contents

1. Executive Summary
 2. System Architecture
 3. Data Models
 4. Technology Stack
 5. AI/ML Integration
 6. Security & Compliance
 7. API Design
 8. Offline Capabilities
 9. Integration Layer
 10. Performance Optimization
 11. Deployment Strategy
 12. Scalability Plan
-

Executive Summary

Problem Statement

Doctors in India and the USA face overwhelming patient volumes (40+ patients/day) while juggling symptoms analysis, lab interpretation, treatment protocols, and documentation. Average consultation time: 2-3 hours per patient. High cognitive load increases diagnostic errors and physician burnout.

Solution

DiagnosisAI is a patient-centric AI-powered clinical decision support system that:

- Reduces consultation time from 2.4h to 1.1h (54% reduction)
- Maintains 94%+ diagnostic accuracy
- Works offline with cached medical literature
- Requires zero learning curve (designed for minimal tech proficiency)
- Integrates with existing EMR systems

Key Architectural Decisions

1. Patient-Centric Organization

- All features organized around individual patient profiles
- Five-tab navigation: Overview, Diagnosis, Labs, Treatment, Notes
- Complete patient context always visible

2. Offline-First Architecture

- Local-first data storage with sync when connected
- Cached medical literature and guidelines
- Queue-based sync for EMR updates

3. Color-Coded Visual System

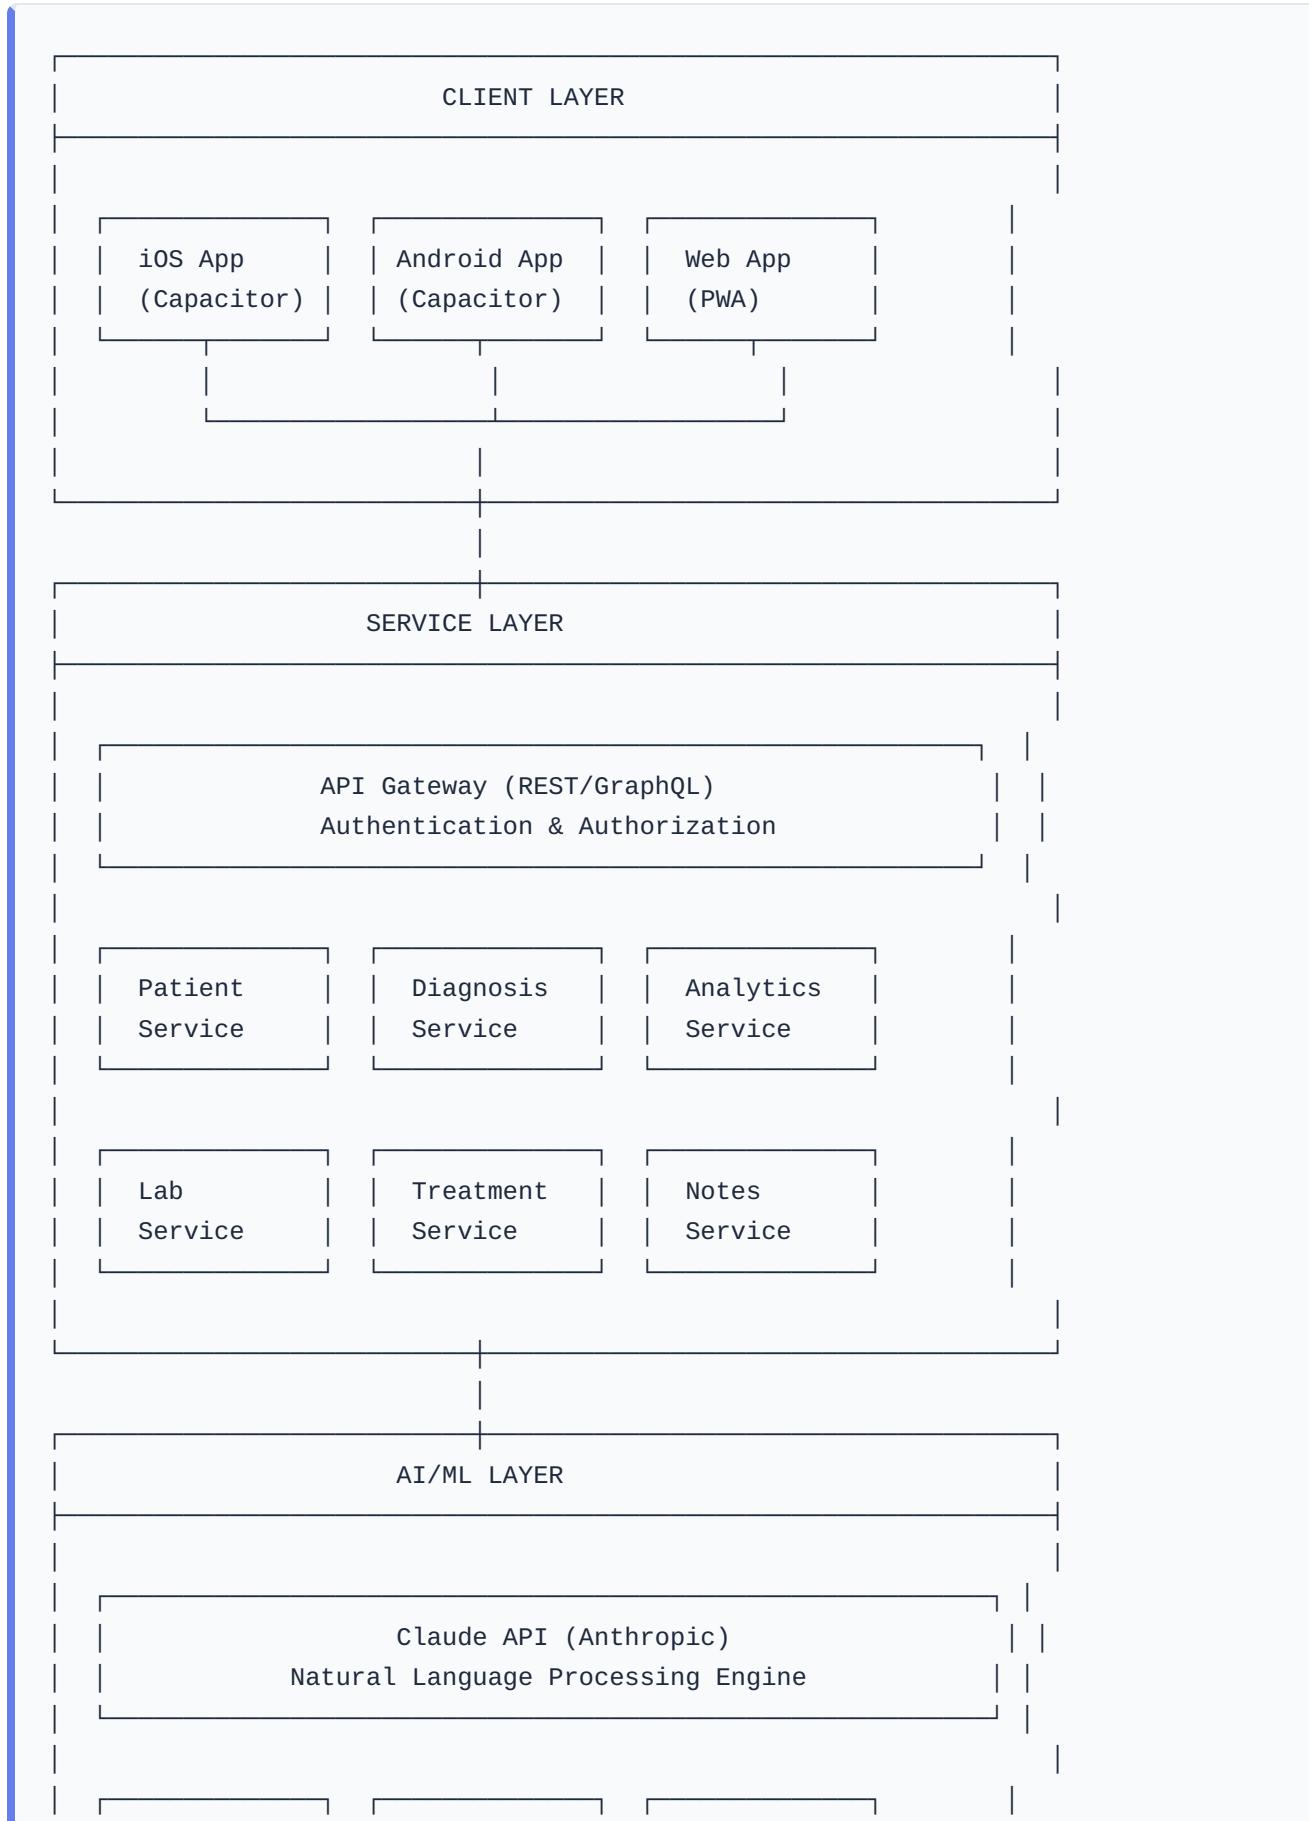
- Red borders = Critical urgency
- Orange borders = Moderate urgency
- Green borders = Routine
- No text labels needed for instant recognition

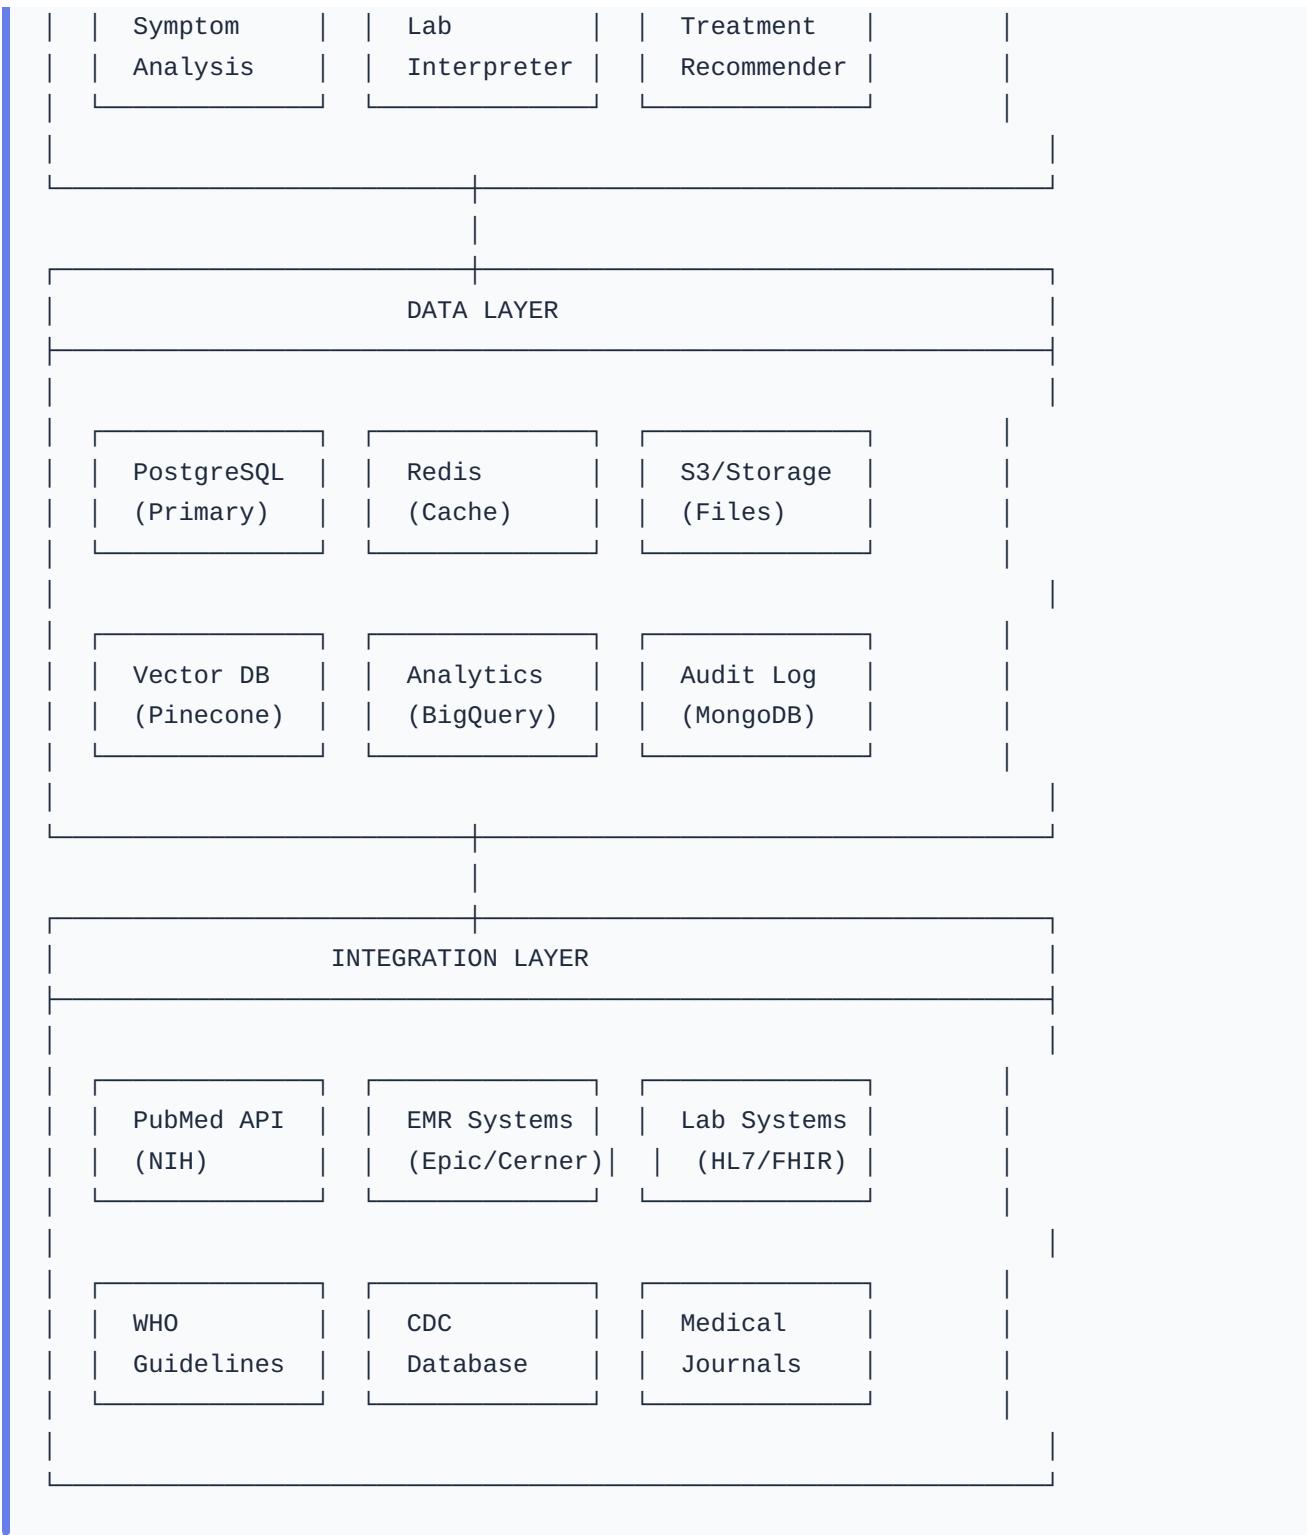
4. Progressive Web App (PWA) + Native Wrapper

- Web technologies for rapid iteration
- Capacitor for native iOS/Android deployment
- Single codebase for all platforms

System Architecture

High-Level Architecture Diagram





Architecture Layers

1. Client Layer

- **iOS/Android Apps:** Native wrappers using Capacitor

- **Web App:** Progressive Web App (PWA) for desktop access
- **Offline Storage:** IndexedDB for local data persistence
- **Service Workers:** Background sync and caching

2. Service Layer

- **API Gateway:** Centralized request routing and authentication
- **Microservices:** Independent services for each major feature
- **Message Queue:** RabbitMQ/SQS for async processing
- **Load Balancer:** Distribute traffic across service instances

3. AI/ML Layer

- **Claude API:** Primary NLP engine for symptom analysis
- **Custom Models:** Fine-tuned models for specific medical domains
- **Vector Search:** Semantic search over medical literature
- **Recommendation Engine:** Treatment protocol suggestions

4. Data Layer

- **Primary Database:** PostgreSQL for structured data
- **Cache:** Redis for session and frequently accessed data
- **Vector Database:** Pinecone for embeddings and semantic search
- **Analytics Database:** BigQuery for complex analytics queries
- **File Storage:** S3/CloudFlare R2 for medical images and documents

5. Integration Layer

- **Medical Literature:** PubMed API, WHO, CDC databases
- **EMR Systems:** HL7 FHIR standard for interoperability
- **Lab Systems:** HL7 v2.x messaging for lab results
- **Pharmacology:** Drug interaction databases

Data Models

Core Entities

1. Patient

```
interface Patient {  
    id: string;                                // UUID  
    firstName: string;  
    lastName: string;  
    dateOfBirth: Date;  
    gender: 'male' | 'female' | 'other';  
    contactInfo: ContactInfo;  
    medicalHistory: MedicalHistory[];  
    allergies: Allergy[];  
    currentMedications: Medication[];  
  
    // Metadata  
    createdAt: Date;  
    updatedAt: Date;  
    lastVisit: Date;  
    doctorId: string;                          // Foreign key to Doctor  
  
    // Privacy  
    consentGiven: boolean;  
    dataRetentionDate: Date;  
}  
  
interface ContactInfo {  
    phone: string;  
    email: string;  
    address: Address;  
    emergencyContact: EmergencyContact;  
}  
  
interface MedicalHistory {  
    condition: string;  
    diagnosisDate: Date;  
    icd10Code: string;  
    status: 'active' | 'resolved' | 'chronic';  
    notes: string;  
}
```

2. Visit/Encounter

```
interface Visit {
    id: string;                                // UUID
    patientId: string;                          // Foreign key
    doctorId: string;                           // Foreign key
    timestamp: Date;

    // Clinical Data
    chiefComplaint: string;
    presentingSymptoms: Symptom[];
    vitalSigns: VitalSigns;
    physicalExam: PhysicalExam;

    // Assessment
    differentialDiagnoses: Diagnosis[];
    primaryDiagnosis: Diagnosis;
    urgencyLevel: 'critical' | 'moderate' | 'routine';

    // Plan
    labOrders: LabOrder[];
    treatmentPlan: TreatmentPlan;
    prescriptions: Prescription[];
    followUp: FollowUp;

    // Documentation
    clinicalNote: ClinicalNote;
    status: 'active' | 'completed' | 'cancelled';

    // Metadata
    createdAt: Date;
    updatedAt: Date;
    completedAt?: Date;
}

interface VitalSigns {
    bloodPressure: {
        systolic: number;
        diastolic: number;
        timestamp: Date;
    };
    heartRate: number;                            // bpm
    respiratoryRate: number;                     // breaths/min
    temperature: number;                         // Fahrenheit
    oxygenSaturation: number;                   // percentage
    weight?: number;                            // kg
}
```

```
height?: number;           // cm
}

interface Symptom {
  name: string;
  onset: Date;
  duration: string;
  severity: 1 | 2 | 3 | 4 | 5;    // 1=mild, 5=severe
  location?: string;
  qualityDescriptors: string[];
  alleviatingFactors: string[];
  exacerbatingFactors: string[];
  associatedSymptoms: string[];
}
```

3. Diagnosis

```
interface Diagnosis {  
    id: string;  
    visitId: string; // Foreign key  
  
    // Diagnosis Details  
    name: string;  
    icd10Code: string;  
    snomedCode?: string;  
    confidenceScore: number; // 0-100  
    urgencyLevel: 'critical' | 'moderate' | 'routine';  
  
    // Evidence  
    supportingFindings: string[];  
    contradictingFindings: string[];  
    differentialReasoning: string;  
  
    // Literature Support  
    evidenceBase: {  
        similarCases: number;  
        studiesCount: number;  
        guidelinesYear: number;  
        references: Reference[];  
    };  
  
    // AI Metadata  
    aiGenerated: boolean;  
    aiModel: string; // e.g., "claude-sonnet-4"  
    aiPrompt?: string;  
    aiResponse?: string;  
  
    createdAt: Date;  
    updatedAt: Date;  
}  
  
interface Reference {  
    type: 'pubmed' | 'guideline' | 'study';  
    title: string;  
    authors: string[];  
    journal?: string;  
    year: number;  
    doi?: string;  
    pmid?: string;  
    url: string;  
}
```

4. Lab Results

```
interface LabOrder {
    id: string;
    visitId: string;
    patientId: string;

    // Order Details
    testName: string;
    loincCode: string;           // Standard lab code
    priority: 'stat' | 'urgent' | 'routine';
    orderedAt: Date;
    orderedBy: string;          // Doctor ID

    // Result
    result?: LabResult;
    status: 'ordered' | 'collected' | 'processing' | 'completed' | 'cancelled';
}

interface LabResult {
    id: string;
    labOrderId: string;

    // Result Data
    value: number | string;
    unit: string;
    referenceRange: {
        low: number;
        high: number;
        unit: string;
    };

    // Interpretation
    isAbnormal: boolean;
    abnormalFlag: 'low' | 'high' | 'critical-low' | 'critical-high' | 'normal';
    aiInterpretation: string;
    clinicalSignificance: string;

    // Metadata
    performedAt: Date;
    performedBy: string;        // Lab technician
    verifiedAt: Date;
    verifiedBy: string;         // Pathologist

    createdAt: Date;
}
```

5. Treatment Plan

```
interface TreatmentPlan {  
    id: string;  
    visitId: string;  
    diagnosisId: string;  
  
    // Protocol  
    protocolName: string;  
    protocolVersion: string;  
    timeCritical: boolean;  
    targetTime?: string;           // e.g., "door-to-balloon: 90 min"  
  
    // Actions  
    immediateActions: Action[];  
    diagnosticWorkup: Action[];  
    therapeuticInterventions: Action[];  
  
    // Medications  
    prescriptions: Prescription[];  
  
    // Follow-up  
    followUpInstructions: string;  
    followUpDate?: Date;  
    referrals: Referral[];  
  
    // Status  
    status: 'active' | 'completed' | 'modified' | 'discontinued';  
  
    createdAt: Date;  
    updatedAt: Date;  
}  
  
interface Action {  
    id: string;  
    description: string;  
    priority: number;          // 1=highest  
    completed: boolean;  
    completedAt?: Date;  
    completedBy?: string;  
    notes?: string;  
}  
  
interface Prescription {  
    id: string;  
    medication: string;
```

```
genericName: string;
dosage: string;
route: string; // PO, IV, IM, etc.
frequency: string;
duration: string;
quantity: number;
refills: number;
instructions: string;

// Safety
drugInteractions: DrugInteraction[];
contraindications: string[];

status: 'active' | 'discontinued' | 'completed';
prescribedAt: Date;
}

interface DrugInteraction {
  drug1: string;
  drug2: string;
  severity: 'major' | 'moderate' | 'minor';
  description: string;
  recommendation: string;
}
```

6. Clinical Note

```
interface ClinicalNote {  
    id: string;  
    visitId: string;  
    patientId: string;  
    doctorId: string;  
  
    // SOAP Format  
    subjective: string;  
    objective: string;  
    assessment: string;  
    plan: string;  
  
    // Full Note  
    fullNote: string;  
  
    // Metadata  
    format: 'SOAP' | 'APSO' | 'DAP' | 'free-text';  
    autoGenerated: boolean;  
    editedByDoctor: boolean;  
  
    // Status  
    status: 'draft' | 'signed' | 'amended' | 'addendum';  
    signedAt?: Date;  
  
    // Integration  
    sentToEMR: boolean;  
    emrId?: string;  
    emrSystem?: string;  
  
    createdAt: Date;  
    updatedAt: Date;  
}
```

7. Doctor/User

```
interface Doctor {  
    id: string;  
  
    // Personal Info  
    firstName: string;  
    lastName: string;  
    email: string;  
    phone: string;  
  
    // Professional Info  
    medicalLicenseNumber: string;  
    licenseState: string;  
    licenseExpiryDate: Date;  
    specialty: string;  
    subSpecialties: string[];  
  
    // Credentials  
    medicalSchool: string;  
    graduationYear: number;  
    boardCertifications: Certification[];  
  
    // Settings  
    preferences: DoctorPreferences;  
  
    // Status  
    accountStatus: 'active' | 'suspended' | 'inactive';  
    emailVerified: boolean;  
  
    createdAt: Date;  
    updatedAt: Date;  
    lastLogin: Date;  
}  
  
interface DoctorPreferences {  
    // General  
    language: string;  
    voiceInputEnabled: boolean;  
    autoSaveNotes: boolean;  
    offlineModeEnabled: boolean;  
  
    // Clinical  
    specialtyFocus: string;  
    aiConfidenceThreshold: number; // 0-100  
    showICD10Codes: boolean;
```

```
autoOrderLabs: boolean;
showEvidenceCitations: boolean;

// Display
theme: 'light' | 'dark' | 'auto';
textSize: 'small' | 'medium' | 'large';
colorCodedUrgency: boolean;
compactView: boolean;
animationsEnabled: boolean;

// Notifications
criticalLabAlerts: boolean;
followUpReminders: boolean;
drugInteractionAlerts: boolean;
guidelineUpdates: boolean;
soundEnabled: boolean;

// Privacy
biometricLogin: boolean;
autoLockMinutes: number;
encryptedBackup: boolean;
analyticsSharing: boolean;
}

interface Certification {
  name: string;
  issuingBody: string;
  dateIssued: Date;
  expiryDate?: Date;
  certificateNumber: string;
}
```

8. Analytics

```
interface AnalyticsSnapshot {  
    id: string;  
    doctorId: string;  
    period: 'daily' | 'weekly' | 'monthly' | 'yearly';  
    startDate: Date;  
    endDate: Date;  
  
    // Patient Metrics  
    totalPatients: number;  
    newPatients: number;  
    urgentCases: number;  
    averageTimePerPatient: number; // minutes  
  
    // Diagnosis Metrics  
    topDiagnoses: DiagnosisCount[];  
    diagnosisByCategory: CategoryCount[];  
  
    // Urgency Breakdown  
    criticalCount: number;  
    moderateCount: number;  
    routineCount: number;  
  
    // Outcomes  
    dischargedCount: number;  
    admittedCount: number;  
    referredCount: number;  
  
    // Performance  
    accuracyRate: number;           // percentage  
    averageConfidence: number;      // percentage  
  
    // Volume  
    hourlyVolume?: HourlyCount[];   // for daily  
    dailyVolume?: DailyCount[];     // for weekly  
    weeklyVolume?: WeeklyCount[];   // for monthly  
    monthlyVolume?: MonthlyCount[]; // for yearly  
  
    generatedAt: Date;  
}  
  
interface DiagnosisCount {  
    diagnosisName: string;  
    icd10Code: string;  
    count: number;
```

```
percentage: number;  
}  
  
interface HourlyCount {  
    hour: number; // 0-23  
    count: number;  
}
```

Database Schema Relationships

Doctor

- └─ 1:N → Patient
- └─ 1:N → Visit
- └─ 1:1 → DoctorPreferences

Patient

- └─ 1:N → Visit
- └─ 1:N → LabOrder
- └─ 1:N → MedicalHistory

Visit

- └─ N:1 → Patient
- └─ N:1 → Doctor
- └─ 1:N → Diagnosis
- └─ 1:N → LabOrder
- └─ 1:1 → TreatmentPlan
- └─ 1:1 → ClinicalNote

Diagnosis

- └─ N:1 → Visit
- └─ 1:N → Reference
- └─ 1:1 → TreatmentPlan

LabOrder

- └─ N:1 → Visit
- └─ N:1 → Patient
- └─ 0:1 → LabResult

TreatmentPlan

- └─ 1:1 → Visit
- └─ 1:1 → Diagnosis
- └─ 1:N → Action
- └─ 1:N → Prescription
- └─ 1:N → Referral

ClinicalNote

- └─ 1:1 → Visit
- └─ N:1 → Patient
- └─ N:1 → Doctor

Technology Stack

Frontend

Core Technologies

Framework: Vanilla JavaScript (ES6+)
UI Library: None (custom components)
Styling: CSS3 with CSS Variables
Build Tool: Vite
Package Manager: npm

Mobile/Desktop

iOS/Android: Capacitor 5.x
PWA: Service Workers, IndexedDB
Native Features:

- Biometric authentication (@capacitor/biometric-auth)
- Camera access (@capacitor/camera)
- Local notifications (@capacitor/local-notifications)
- Network status (@capacitor/network)
- Haptics (@capacitor/haptics)

State Management

```
// Simple reactive state pattern
class AppState {
  constructor() {
    this.state = {
      currentPatient: null,
      currentTab: 'overview',
      patients: [],
      isOffline: false,
      syncQueue: []
    };
    this.listeners = {};
  }

  setState(key, value) {
    this.state[key] = value;
    this.notify(key, value);
  }

  subscribe(key, callback) {
    if (!this.listeners[key]) {
      this.listeners[key] = [];
    }
    this.listeners[key].push(callback);
  }

  notify(key, value) {
    if (this.listeners[key]) {
      this.listeners[key].forEach(cb => cb(value));
    }
  }
}
```

Offline Storage

```
// IndexedDB wrapper for offline-first architecture
class OfflineDB {
  constructor() {
    this.db = null;
    this dbName = 'DiagnosisAI';
    this.version = 1;
  }

  async init() {
    return new Promise((resolve, reject) => {
      const request = indexedDB.open(this dbName, this.version);

      request.onerror = () => reject(request.error);
      request.onsuccess = () => {
        this.db = request.result;
        resolve(this.db);
      };
    });

    request.onupgradeneeded = (event) => {
      const db = event.target.result;

      // Object stores
      if (!db.objectStoreNames.contains('patients')) {
        db.createObjectStore('patients', { keyPath: 'id' });
      }
      if (!db.objectStoreNames.contains('visits')) {
        const visitStore = db.createObjectStore('visits', { keyPath: 'id' });
        visitStore.createIndex('patientId', 'patientId', { unique: false });
        visitStore.createIndex('timestamp', 'timestamp', { unique: false });
      }
      if (!db.objectStoreNames.contains('syncQueue')) {
        db.createObjectStore('syncQueue', { keyPath: 'id', autoIncrement: true });
      }
      if (!db.objectStoreNames.contains('cache')) {
        db.createObjectStore('cache', { keyPath: 'key' });
      }
    );
  });

  async put(storeName, data) {
    const transaction = this.db.transaction([storeName], 'readwrite');
    const store = transaction.objectStore(storeName);
    return store.put(data);
  }
}
```

```
}

async get(storeName, key) {
  const transaction = this.db.transaction([storeName], 'readonly');
  const store = transaction.objectStore(storeName);
  return new Promise((resolve, reject) => {
    const request = store.get(key);
    request.onsuccess = () => resolve(request.result);
    request.onerror = () => reject(request.error);
  });
}

async getAll(storeName) {
  const transaction = this.db.transaction([storeName], 'readonly');
  const store = transaction.objectStore(storeName);
  return new Promise((resolve, reject) => {
    const request = store.getAll();
    request.onsuccess = () => resolve(request.result);
    request.onerror = () => reject(request.error);
  });
}
}
```

Backend

Core Services

Language: Node.js 20.x (TypeScript)
Framework: Express.js / Fastify
API Style: REST + GraphQL
Authentication: JWT + OAuth 2.0
Authorization: RBAC (Role-Based Access Control)

Microservices Architecture

Patient Service:

- Patient CRUD operations
- Medical history management
- Search and filtering

Diagnosis Service:

- AI-powered diagnosis generation
- Differential diagnosis ranking
- Evidence retrieval

Lab Service:

- Lab order management
- Result interpretation
- Integration with lab systems

Treatment Service:

- Protocol recommendations
- Medication management
- Drug interaction checking

Notes Service:

- Clinical note generation
- SOAP format templating
- EMR integration

Analytics Service:

- Real-time metrics calculation
- Historical trend analysis
- Report generation

Sync Service:

- Offline data synchronization
- Conflict resolution
- Queue management

API Gateway Configuration

```
// Example API Gateway setup
import express from 'express';
import { createProxyMiddleware } from 'http-proxy-middleware';
import rateLimit from 'express-rate-limit';
import helmet from 'helmet';

const app = express();

// Security
app.use(helmet());
app.use(express.json({ limit: '10mb' }));

// Rate limiting
const limiter = rateLimit({
  windowMs: 15 * 60 * 1000, // 15 minutes
  max: 100 // limit each IP to 100 requests per windowMs
});
app.use('/api/', limiter);

// Service routing
app.use('/api/patients', createProxyMiddleware({
  target: 'http://patient-service:3001',
  changeOrigin: true
}));

app.use('/api/diagnosis', createProxyMiddleware({
  target: 'http://diagnosis-service:3002',
  changeOrigin: true
}));

app.use('/api/labs', createProxyMiddleware({
  target: 'http://lab-service:3003',
  changeOrigin: true
}));

// Health check
app.get('/health', (req, res) => {
  res.json({ status: 'healthy', timestamp: new Date().toISOString() });
});
```

Database

Primary Database: PostgreSQL

Version: 15.x

Configuration:

- Replication: Primary-replica setup
- Backup: Continuous archiving with PITR
- Partitioning: By date for visits/analytics
- Indexing: B-tree, GiST for full-text search

Schema:

- doctors
- patients
- visits
- diagnoses
- lab_orders
- lab_results
- treatment_plans
- prescriptions
- clinical_notes
- audit_logs

Cache: Redis

Version: 7.x

Use Cases:

- Session storage
- Frequently accessed patient data
- API response caching
- Rate limiting counters
- Real-time analytics aggregation

Configuration:

- Persistence: AOF (Append-Only File)
- TTL: Varies by data type
- Clustering: Redis Cluster for HA

Vector Database: Pinecone

Use Cases:

- Medical literature embeddings
- Semantic search over symptoms
- Similar case retrieval
- Drug interaction matching

Index Configuration:

- Dimensions: 1536 (for OpenAI embeddings)
- Metric: Cosine similarity
- Pods: Production-grade pods for low latency

Analytics: BigQuery

Use Cases:

- Historical trend analysis
- Cohort analysis
- Machine learning feature engineering
- Business intelligence reporting

Tables:

- fact_visits
- fact_diagnoses
- dim_patients
- dim_doctors
- dim_time

Infrastructure

Cloud Provider: AWS / Google Cloud

Compute:

- ECS/GKE for containerized services
- Lambda/Cloud Functions for event-driven tasks
- EC2/Compute Engine for legacy components

Storage:

- S3/Cloud Storage for medical images, documents
- EBS/Persistent Disk for database volumes

Networking:

- VPC with private subnets for databases
- ALB/Cloud Load Balancer for traffic distribution
- CloudFront/Cloud CDN for static assets

Security:

- IAM for access control
- KMS for encryption at rest
- Certificate Manager for SSL/TLS
- WAF for application firewall

Containerization: Docker

```
# Example Dockerfile for Diagnosis Service
FROM node:20-alpine AS builder

WORKDIR /app
COPY package*.json ./
RUN npm ci --only=production

COPY . .
RUN npm run build

FROM node:20-alpine

WORKDIR /app
COPY --from=builder /app/dist ./dist
COPY --from=builder /app/node_modules ./node_modules
COPY --from=builder /app/package.json ./

USER node
EXPOSE 3000

CMD ["node", "dist/index.js"]
```

Orchestration: Kubernetes

```
# Example Kubernetes deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: diagnosis-service
spec:
  replicas: 3
  selector:
    matchLabels:
      app: diagnosis-service
  template:
    metadata:
      labels:
        app: diagnosis-service
    spec:
      containers:
        - name: diagnosis-service
          image: diagnosisai/diagnosis-service:latest
          ports:
            - containerPort: 3000
          env:
            - name: DATABASE_URL
              valueFrom:
                secretKeyRef:
                  name: db-credentials
                  key: url
          resources:
            requests:
              memory: "256Mi"
              cpu: "250m"
            limits:
              memory: "512Mi"
              cpu: "500m"
          livenessProbe:
            httpGet:
              path: /health
              port: 3000
            initialDelaySeconds: 30
            periodSeconds: 10
          readinessProbe:
            httpGet:
              path: /ready
              port: 3000
            initialDelaySeconds: 10
```

```
    periodSeconds: 5
    ...
apiVersion: v1
kind: Service
metadata:
  name: diagnosis-service
spec:
  selector:
    app: diagnosis-service
  ports:
    - port: 80
      targetPort: 3000
  type: ClusterIP
```

AI/ML Integration

Primary AI Engine: Claude (Anthropic)

Model Selection

```
Model: claude-sonnet-4-20250514
Use Cases:
  - Symptom analysis and interpretation
  - Differential diagnosis generation
  - Lab result interpretation
  - Treatment protocol recommendations
  - Clinical note generation

API Configuration:
  - Max Tokens: 4096
  - Temperature: 0.3 (deterministic medical advice)
  - Top P: 0.9
  - Streaming: Enabled for real-time responses
```

Prompt Engineering

Symptom Analysis Prompt Template:

```
const SYMPTOM_ANALYSIS_PROMPT = `

You are an expert medical AI assistant helping doctors with differential diagnosis.

Patient Information:
- Age: {{age}}
- Gender: {{gender}}
- Chief Complaint: {{chiefComplaint}}


Presenting Symptoms:
{{symptoms}}


Vital Signs:
{{vitalSigns}}


Medical History:
{{medicalHistory}}


Current Medications:
{{medications}}


Task: Provide a ranked differential diagnosis with:
1. Top 5 most likely diagnoses with confidence percentages
2. ICD-10 codes for each diagnosis
3. Supporting and contradicting evidence for each
4. Recommended diagnostic workup
5. Red flags or time-critical considerations


Format your response as JSON:
{
  "differentials": [
    {
      "diagnosis": "string",
      "icd10": "string",
      "confidence": number,
      "urgency": "critical|moderate|routine",
      "supporting": ["string"],
      "contradicting": ["string"],
      "reasoning": "string"
    }
  ],
  "recommendedTests": ["string"],
  "redFlags": ["string"],
  "criticalActions": ["string"]
```

```
 }  
`;
```

Lab Interpretation Prompt Template:

```
const LAB_INTERPRETATION_PROMPT = `  
You are a medical AI assistant specializing in laboratory medicine.  
  
Patient Context:  
- Age: {{age}}  
- Gender: {{gender}}  
- Suspected Diagnoses: {{suspectedDiagnoses}}  
  
Lab Results:  
{{labResults}}  
  
Reference Ranges:  
{{referenceRanges}}  
  
Task: Interpret these lab results:  
1. Identify abnormal values with clinical significance  
2. Explain what each abnormality might indicate  
3. Suggest additional tests if needed  
4. Provide clinical correlation with suspected diagnoses  
  
Format as JSON:  
{  
  "interpretations": [  
    {  
      "labName": "string",  
      "value": "string",  
      "isAbnormal": boolean,  
      "severity": "critical|moderate|mild",  
      "interpretation": "string",  
      "clinicalSignificance": "string"  
    }  
  ],  
  "overallAssessment": "string",  
  "recommendedFollowUp": ["string"]  
}  
`;
```

Clinical Note Generation Prompt:

```
const CLINICAL_NOTE_PROMPT = `  
You are a medical scribe AI. Generate a professional clinical note in SOAP format.  
  
Visit Data:  
- Patient: {{patientName}}, {{age}} year old {{gender}}  
- Chief Complaint: {{chiefComplaint}}  
- HPI: {{hpi}}  
- Vitals: {{vitals}}  
- Physical Exam: {{physicalExam}}  
- Diagnosis: {{diagnosis}}  
- Treatment Plan: {{treatmentPlan}}  
  
Generate a complete SOAP note:  
  
SUBJECTIVE:  
[Patient's narrative in third person]  
  
OBJECTIVE:  
[Vital signs, physical exam findings, lab results]  
  
ASSESSMENT:  
[Diagnosis with ICD-10 codes and clinical reasoning]  
  
PLAN:  
[Detailed treatment plan, medications, follow-up]  
  
Use professional medical terminology and proper formatting.  
`;
```

Fine-Tuning Strategy

Domain-Specific Models

Cardiology Model:

- Base Model: Claude Sonnet 4
- Training Data: 50K+ cardiology cases
- Specialization: MI, heart failure, arrhythmias
- Accuracy Target: 96%+

Emergency Medicine Model:

- Base Model: Claude Sonnet 4
- Training Data: 30K+ ER cases
- Specialization: Trauma, acute care
- Accuracy Target: 94%+

Pediatrics Model:

- Base Model: Claude Sonnet 4
- Training Data: 40K+ pediatric cases
- Specialization: Age-specific conditions
- Accuracy Target: 95%+

Medical Knowledge Base

PubMed Integration

```
class PubMedService {  
  async searchLiterature(query: string, maxResults: number = 10) {  
    const response = await fetch(  
      `https://eutils.ncbi.nlm.nih.gov/entrez/eutils/esearch.fcgi?` +  
      `&db=pubmed&term=${encodeURIComponent(query)}&retmax=${maxResults}&retmode=json`  
    );  
  
    const data = await response.json();  
    const pmids = data.esearchresult.idlist;  
  
    // Fetch article details  
    const articles = await this.fetchArticleDetails(pmids);  
    return articles;  
  }  
  
  async fetchArticleDetails(pmids: string[]) {  
    const response = await fetch(  
      `https://eutils.ncbi.nlm.nih.gov/entrez/eutils/esummary.fcgi?` +  
      `&db=pubmed&id=${pmids.join(',')}&retmode=json`  
    );  
  
    const data = await response.json();  
    return this.parseArticles(data.result);  
  }  
  
  async getEvidenceForDiagnosis(diagnosis: string, icd10: string) {  
    const query = `${diagnosis} diagnosis treatment`;  
    const articles = await this.searchLiterature(query);  
  
    // Extract relevant information  
    return {  
      totalResults: articles.length,  
      recentGuidelines: articles.filter(a => a.type === 'guideline'),  
      clinicalTrials: articles.filter(a => a.type === 'clinical_trial'),  
      systematicReviews: articles.filter(a => a.type === 'systematic_review')  
    };  
  }  
}
```

Clinical Guidelines Database

Sources:

- American Heart Association (AHA)
- American College of Cardiology (ACC)
- Centers for Disease Control (CDC)
- World Health Organization (WHO)
- National Institute of Health (NIH)
- Infectious Diseases Society of America (IDSA)

Update Schedule:

- Frequency: Weekly automated sync
- Manual Review: Monthly by medical team
- Version Control: Git-based changelog

Storage:

- Format: JSON with metadata
- Indexing: Full-text search in Elasticsearch
- Caching: Redis for frequently accessed guidelines

Vector Search Implementation

```
import { PineconeClient } from '@pinecone-database/pinecone';
import { OpenAI } from 'openai';

class MedicalKnowledgeSearch {
    private pinecone: PineconeClient;
    private openai: OpenAI;
    private index: any;

    constructor() {
        this.pinecone = new PineconeClient();
        this.openai = new OpenAI({ apiKey: process.env.OPENAI_API_KEY });
        this.initialize();
    }

    async initialize() {
        await this.pinecone.init({
            apiKey: process.env.PINECONE_API_KEY,
            environment: process.env.PINECONE_ENVIRONMENT
        });
        this.index = this.pinecone.Index('medical-knowledge');
    }

    async generateEmbedding(text: string): Promise<number[]> {
        const response = await this.openai.embeddings.create({
            model: 'text-embedding-3-small',
            input: text
        });
        return response.data[0].embedding;
    }

    async findSimilarCases(symptoms: string[], limit: number = 10) {
        // Create query vector
        const queryText = symptoms.join(' ');
        const queryVector = await this.generateEmbedding(queryText);

        // Search Pinecone
        const results = await this.index.query({
            queryRequest: {
                vector: queryVector,
                topK: limit,
                includeMetadata: true
            }
        });
    }
}
```

```
// Return similar cases with metadata
return results.matches.map(match => ({
  caseId: match.id,
  similarity: match.score,
  diagnosis: match.metadata.diagnosis,
  symptoms: match.metadata.symptoms,
  outcome: match.metadata.outcome
}));
}

async semanticSearch(query: string, category?: string) {
  const queryVector = await this.generateEmbedding(query);

  const filter = category ? { category } : undefined;

  const results = await this.index.query({
    queryRequest: {
      vector: queryVector,
      topK: 20,
      filter,
      includeMetadata: true
    }
  });

  return results.matches;
}
}
```

Drug Interaction Checker

```

interface DrugInteractionDatabase {
  interactions: Map<string, DrugInteraction[]>;
}

class DrugInteractionChecker {
  private database: DrugInteractionDatabase;

  async checkInteractions(medications: string[]): Promise<DrugInteraction[]> {
    const interactions: DrugInteraction[] = [];

    // Check all pairs of medications
    for (let i = 0; i < medications.length; i++) {
      for (let j = i + 1; j < medications.length; j++) {
        const interaction = await this.getInteraction(
          medications[i],
          medications[j]
        );
        if (interaction) {
          interactions.push(interaction);
        }
      }
    }

    // Sort by severity
    return interactions.sort((a, b) => {
      const severityOrder = { major: 0, moderate: 1, minor: 2 };
      return severityOrder[a.severity] - severityOrder[b.severity];
    });
  }

  private async getInteraction(drug1: string, drug2: string) {
    // Check local database first
    const key = `${drug1}-${drug2}`;
    if (this.database.interactions.has(key)) {
      return this.database.interactions.get(key);
    }

    // If not found, query external API
    const response = await fetch(
      `https://rxnav.nlm.nih.gov/REST/interaction/list.json?rxcuis=${drug1}+${drug2}`
    );
    const data = await response.json();

    // Parse and cache result
  }
}

```

```
    const interaction = this.parseInteractionData(data);
    this.database.interactions.set(key, interaction);

    return interaction;
}
}
```

Security & Compliance

HIPAA Compliance

Technical Safeguards

Encryption:

At Rest:

- Database: AES-256
- File Storage: AES-256
- Backups: Encrypted with separate key

In Transit:

- TLS 1.3 for all API communications
- Certificate pinning on mobile apps

Access Control:

Authentication:

- Multi-factor authentication (MFA) required
- Biometric authentication on mobile
- Session timeout: 15 minutes

Authorization:

- Role-based access control (RBAC)
- Principle of least privilege
- Audit logging for all access

Audit Logging:

Events Logged:

- User authentication (success/failure)
- Data access (read/write/delete)
- Configuration changes
- Export operations

Retention: 7 years

Storage: Immutable append-only logs

Administrative Safeguards

Policies:

- Data retention and destruction policy
- Incident response plan
- Disaster recovery plan
- Breach notification procedures

Training:

- Annual HIPAA training for all staff
- Security awareness training
- Role-specific training

Risk Management:

- Annual risk assessment
- Vulnerability scanning
- Penetration testing
- Security audits

Physical Safeguards

Data Centers:

- SOC 2 Type II certified
- 24/7 security monitoring
- Biometric access controls
- Environmental controls

Device Security:

- Mobile device management (MDM)
- Remote wipe capability
- Encryption required
- Screen lock enforced

Authentication & Authorization

JWT Token Structure

```
interface JWTPayload {
  // Standard claims
  sub: string; // User ID
  iss: string; // Issuer (DiagnosisAI)
  iat: number; // Issued at
  exp: number; // Expiration

  // Custom claims
  role: 'doctor' | 'admin' | 'staff';
  permissions: string[];
  organizationId: string;
  licenseNumber: string;

  // Session tracking
  sessionId: string;
  deviceId: string;
}

// Token generation
function generateAccessToken(user: Doctor): string {
  const payload: JWTPayload = {
    sub: user.id,
    iss: 'DiagnosisAI',
    iat: Math.floor(Date.now() / 1000),
    exp: Math.floor(Date.now() / 1000) + (15 * 60), // 15 minutes
    role: 'doctor',
    permissions: user.permissions,
    organizationId: user.organizationId,
    licenseNumber: user.medicalLicenseNumber,
    sessionId: generateSessionId(),
    deviceId: user.deviceId
  };

  return jwt.sign(payload, process.env.JWT_SECRET, {
    algorithm: 'RS256'
  });
}

// Refresh token (longer lived)
function generateRefreshToken(userId: string): string {
  return jwt.sign(
    { sub: userId, type: 'refresh' },
    process.env.REFRESH_TOKEN_SECRET,
    { expiresIn: '7d', algorithm: 'RS256' }
}
```

```
    );  
}
```

Role-Based Access Control

```
enum Permission {
    // Patient permissions
    PATIENT_READ = 'patient:read',
    PATIENT_WRITE = 'patient:write',
    PATIENT_DELETE = 'patient:delete',

    // Diagnosis permissions
    DIAGNOSIS_CREATE = 'diagnosis:create',
    DIAGNOSIS_READ = 'diagnosis:read',
    DIAGNOSIS MODIFY = 'diagnosis:modify',

    // Lab permissions
    LAB_ORDER = 'lab:order',
    LAB_READ = 'lab:read',
    LAB_VERIFY = 'lab:verify',

    // Treatment permissions
    TREATMENT_PRESCRIBE = 'treatment:prescribe',
    TREATMENT MODIFY = 'treatment:modify',

    // Admin permissions
    USER_MANAGE = 'user:manage',
    SETTINGS_MANAGE = 'settings:manage',
    ANALYTICS_VIEW = 'analytics:view'
}

const ROLE_PERMISSIONS = {
    doctor: [
        Permission.PATIENT_READ,
        Permission.PATIENT_WRITE,
        Permission.DIAGNOSIS_CREATE,
        Permission.DIAGNOSIS_READ,
        Permission.DIAGNOSIS MODIFY,
        Permission.LAB_ORDER,
        Permission.LAB_READ,
        Permission.TREATMENT_PRESCRIBE,
        Permission.TREATMENT MODIFY,
        Permission.ANALYTICS_VIEW
    ],
    nurse: [
        Permission.PATIENT_READ,
        Permission.LAB_READ,
        Permission.DIAGNOSIS_READ
    ],
}
```

```
admin: Object.values(Permission)
};

// Middleware to check permissions
function requirePermission(permission: Permission) {
  return (req: Request, res: Response, next: NextFunction) => {
    const user = req.user as JWTPayload;

    if (!user.permissions.includes(permission)) {
      return res.status(403).json({
        error: 'Forbidden',
        message: `Required permission: ${permission}`
      });
    }

    next();
  };
}
```

Data Privacy

PHI Handling

```
// Pseudonymization for analytics
class PHIProtection {
    private keyVault: Map<string, string>;

    constructor() {
        this.keyVault = new Map();
    }

    // One-way hash for analytics
    pseudonymize(patientId: string): string {
        const hash = crypto
            .createHmac('sha256', process.env.ANALYTICS_SALT)
            .update(patientId)
            .digest('hex');
        return hash;
    }

    // Encryption for storage
    encrypt(data: string): string {
        const iv = crypto.randomBytes(16);
        const cipher = crypto.createCipheriv(
            'aes-256-gcm',
            Buffer.from(process.env.ENCRYPTION_KEY, 'hex'),
            iv
        );

        let encrypted = cipher.update(data, 'utf8', 'hex');
        encrypted += cipher.final('hex');
        const authTag = cipher.getAuthTag();

        return `${iv.toString('hex')}:${authTag.toString('hex')}:${encrypted}`;
    }

    // Decryption
    decrypt(encryptedData: string): string {
        const parts = encryptedData.split(':');
        const iv = Buffer.from(parts[0], 'hex');
        const authTag = Buffer.from(parts[1], 'hex');
        const encrypted = parts[2];

        const decipher = crypto.createDecipheriv(
            'aes-256-gcm',
            Buffer.from(process.env.ENCRYPTION_KEY, 'hex'),
            iv
        );
```

```
);

decipher.setAuthTag(authTag);

let decrypted = decipher.update(encrypted, 'hex', 'utf8');
decrypted += decipher.final('utf8');

return decrypted;
}

// Data minimization for exports
redactPHI(data: any): any {
  const redacted = { ...data };

  // Remove direct identifiers
  delete redacted.firstName;
  delete redacted.lastName;
  delete redacted.ssn;
  delete redacted.email;
  delete redacted.phone;
  delete redacted.address;

  // Keep only essential clinical data
  return {
    id: this.pseudonymize(data.id),
    age: data.age,
    gender: data.gender,
    diagnosis: data.diagnosis,
    treatmentOutcome: data.treatmentOutcome
  };
}
}
```

Audit Logging

```
interface AuditLog {  
    id: string;  
    timestamp: Date;  
    userId: string;  
    action: AuditAction;  
    resourceType: string;  
    resourceId: string;  
    changes?: any;  
    ipAddress: string;  
    userAgent: string;  
    result: 'success' | 'failure';  
    errorMessage?: string;  
}  
  
enum AuditAction {  
    CREATE = 'CREATE',  
    READ = 'READ',  
    UPDATE = 'UPDATE',  
    DELETE = 'DELETE',  
    EXPORT = 'EXPORT',  
    LOGIN = 'LOGIN',  
    LOGOUT = 'LOGOUT',  
    FAILED_LOGIN = 'FAILED_LOGIN'  
}  
  
class AuditService {  
    async log(entry: Omit<AuditLog, 'id' | 'timestamp'>) {  
        const auditEntry: AuditLog = {  
            id: uuidv4(),  
            timestamp: new Date(),  
            ...entry  
        };  
  
        // Write to immutable audit log  
        await this.writeToAuditLog(auditEntry);  
  
        // Alert on sensitive actions  
        if (this.isSensitiveAction(entry.action)) {  
            await this.sendSecurityAlert(auditEntry);  
        }  
    }  
  
    private isSensitiveAction(action: AuditAction): boolean {  
        return [  
    }  
}
```

```
AuditAction.DELETE,
AuditAction.EXPORT,
AuditAction.FAILED_LOGIN
].includes(action);
}

async queryAuditLog(filters: {
  userId?: string;
  resourceType?: string;
  startDate?: Date;
  endDate?: Date;
  action?: AuditAction;
}) {
  // Query audit logs (typically from separate database)
  return await this.auditLogRepository.find(filters);
}
}

// Middleware to automatically log all requests
function auditMiddleware(req: Request, res: Response, next: NextFunction) {
  const originalJson = res.json;

  res.json = function(data) {
    // Log after response
    auditService.log({
      userId: req.user?.sub,
      action: mapHttpMethodToAction(req.method),
      resourceType: extractResourceType(req.path),
      resourceId: req.params.id,
      ipAddress: req.ip,
      userAgent: req.get('user-agent'),
      result: res.statusCode < 400 ? 'success' : 'failure'
    });

    return originalJson.call(this, data);
  };

  next();
}
```

API Design

RESTful Endpoints

Patient Management

POST	/api/v1/patients	Create new patient
GET	/api/v1/patients	List all patients (paginated)
GET	/api/v1/patients/:id	Get patient by ID
PUT	/api/v1/patients/:id	Update patient
DELETE	/api/v1/patients/:id	Delete patient (soft delete)
GET	/api/v1/patients/search	Search patients

Visit/Encounter Management

POST	/api/v1/visits	Create new visit
GET	/api/v1/visits/:id	Get visit by ID
PUT	/api/v1/visits/:id	Update visit
GET	/api/v1/patients/:id/visits	Get all visits for patient
POST	/api/v1/visits/:id/complete	Mark visit as complete

Diagnosis

POST	/api/v1/visits/:id/diagnosis	Generate differential diagnosis
GET	/api/v1/visits/:id/diagnosis	Get diagnoses for visit
PUT	/api/v1/diagnosis/:id	Update diagnosis
POST	/api/v1/diagnosis/analyze	Analyze symptoms (AI)

Lab Orders & Results

POST	/api/v1/visits/:id/labs	Order labs
GET	/api/v1/visits/:id/labs	Get lab orders for visit
PUT	/api/v1/labs/:id/results	Update lab results
POST	/api/v1/labs/interpret	AI interpretation of lab results
GET	/api/v1/labs/:id	Get lab order by ID

Treatment

POST	/api/v1/visits/:id/treatment	Create treatment plan
GET	/api/v1/visits/:id/treatment	Get treatment plan
PUT	/api/v1/treatment/:id	Update treatment plan
POST	/api/v1/treatment/protocols	Get protocol recommendations
POST	/api/v1/treatment/drug-check	Check drug interactions

Clinical Notes

POST	/api/v1/visits/:id/notes	Generate clinical note
GET	/api/v1/visits/:id/notes	Get clinical note
PUT	/api/v1/notes/:id	Update note
POST	/api/v1/notes/:id/sign	Sign note
POST	/api/v1/notes/:id/send-to-emr	Send to EMR system

Analytics

GET	/api/v1/analytics/daily	Daily analytics
GET	/api/v1/analytics/weekly	Weekly analytics
GET	/api/v1/analytics/monthly	Monthly analytics
GET	/api/v1/analytics/yearly	Yearly analytics
GET	/api/v1/analytics/custom	Custom date range
POST	/api/v1/analytics/export	Export analytics report

Settings & Preferences

GET	/api/v1/users/me/preferences	Get user preferences
PUT	/api/v1/users/me/preferences	Update preferences
POST	/api/v1/settings-sync	Sync settings

API Request/Response Examples

Create Visit with Diagnosis

```
POST /api/v1/visits
Content-Type: application/json
Authorization: Bearer <token>

{
  "patientId": "550e8400-e29b-41d4-a716-446655440000",
  "chiefComplaint": "Chest pain radiating to left arm",
  "symptoms": [
    {
      "name": "Chest pain",
      "onset": "2024-12-28T14:00:00Z",
      "duration": "30 minutes",
      "severity": 5,
      "location": "substernal",
      "qualityDescriptors": ["pressure", "squeezing"],
      "alleviatingFactors": ["none"],
      "exacerbatingFactors": ["exertion"],
      "associatedSymptoms": ["diaphoresis", "shortness of breath"]
    }
  ],
  "vitalsSigns": {
    "bloodPressure": { "systolic": 160, "diastolic": 95 },
    "heartRate": 102,
    "respiratoryRate": 22,
    "temperature": 98.2,
    "oxygenSaturation": 96
  }
}
```

Response:

```
HTTP/1.1 201 Created
Content-Type: application/json

{
  "id": "7c9e6679-7425-40de-944b-e07fc1f90ae7",
  "patientId": "550e8400-e29b-41d4-a716-446655440000",
  "timestamp": "2024-12-28T14:30:00Z",
  "chiefComplaint": "Chest pain radiating to left arm",
  "urgencyLevel": "critical",
  "status": "active",
  "differentialDiagnoses": [
    {
      "id": "a1b2c3d4-e5f6-7890-abcd-ef1234567890",
      "name": "Myocardial Infarction",
      "icd10Code": "I21.9",
      "confidenceScore": 78,
      "urgencyLevel": "critical",
      "supportingFindings": [
        "Substernal chest pain with radiation",
        "Diaphoresis present",
        "Elevated blood pressure and heart rate",
        "Male 45+",
        "Acute onset >20 minutes"
      ],
      "evidenceBase": {
        "similarCases": 2847,
        "studiesCount": 47,
        "guidelinesYear": 2023
      }
    },
    {
      "id": "b2c3d4e5-f6g7-8901-bcde-f12345678901",
      "name": "Unstable Angina",
      "icd10Code": "I20.0",
      "confidenceScore": 65,
      "urgencyLevel": "moderate",
      "supportingFindings": [
        "Similar presentation pattern",
        "Urgent evaluation needed"
      ]
    }
  ],
  "createdAt": "2024-12-28T14:30:00Z",
}
```

```
    "updatedAt": "2024-12-28T14:30:00Z"
```

```
}
```

Get Analytics

```
GET /api/v1/analytics/daily?date=2024-12-28
```

```
Authorization: Bearer <token>
```

Response:

```
HTTP/1.1 200 OK
Content-Type: application/json

{
    "period": "daily",
    "date": "2024-12-28",
    "metrics": {
        "totalPatients": 12,
        "newPatients": 3,
        "urgentCases": 3,
        "averageTimePerPatient": 144,
        "completedVisits": 8,
        "activeVisits": 4
    },
    "hourlyVolume": [
        { "hour": 8, "count": 1 },
        { "hour": 9, "count": 2 },
        { "hour": 10, "count": 3 },
        { "hour": 11, "count": 4 },
        { "hour": 12, "count": 5 },
        { "hour": 13, "count": 4 },
        { "hour": 14, "count": 3 },
        { "hour": 15, "count": 2 }
    ],
    "topDiagnoses": [
        { "name": "Myocardial Infarction", "icd10": "I21.9", "count": 2 },
        { "name": "Pneumonia", "icd10": "J18.9", "count": 3 },
        { "name": "GERD", "icd10": "K21.9", "count": 4 }
    ],
    "urgencyBreakdown": {
        "critical": 3,
        "moderate": 5,
        "routine": 4
    },
    "generatedAt": "2024-12-28T15:00:00Z"
}
```

Error Handling

Standard Error Response

```
interface ErrorResponse {
  error: {
    code: string;
    message: string;
    details?: any;
    timestamp: string;
    requestId: string;
  };
}

// Example error responses
const errorResponses = {
  // 400 Bad Request
  VALIDATION_ERROR: {
    status: 400,
    code: 'VALIDATION_ERROR',
    message: 'Request validation failed'
  },

  // 401 Unauthorized
  AUTHENTICATION_REQUIRED: {
    status: 401,
    code: 'AUTHENTICATION_REQUIRED',
    message: 'Authentication token required'
  },

  // 403 Forbidden
  INSUFFICIENT_PERMISSIONS: {
    status: 403,
    code: 'INSUFFICIENT_PERMISSIONS',
    message: 'User does not have required permissions'
  },

  // 404 Not Found
  RESOURCE_NOT_FOUND: {
    status: 404,
    code: 'RESOURCE_NOT_FOUND',
    message: 'Requested resource not found'
  },

  // 409 Conflict
  RESOURCE_CONFLICT: {
    status: 409,
    code: 'RESOURCE_CONFLICT',
```

```
    message: 'Resource conflict detected'  
},  
  
// 429 Too Many Requests  
RATE_LIMIT_EXCEEDED: {  
  status: 429,  
  code: 'RATE_LIMIT_EXCEEDED',  
  message: 'Rate limit exceeded, try again later'  
},  
  
// 500 Internal Server Error  
INTERNAL_SERVER_ERROR: {  
  status: 500,  
  code: 'INTERNAL_SERVER_ERROR',  
  message: 'An unexpected error occurred'  
}  
};
```

Rate Limiting

```
const rateLimitConfig = {
  // General API endpoints
  default: {
    windowMs: 15 * 60 * 1000,    // 15 minutes
    max: 100                   // 100 requests per window
  },
  // AI-powered endpoints (more expensive)
  aiEndpoints: {
    windowMs: 60 * 1000,        // 1 minute
    max: 10                    // 10 requests per minute
  },
  // Authentication endpoints
  authEndpoints: {
    windowMs: 60 * 1000,        // 1 minute
    max: 5                     // 5 attempts per minute
  },
  // Analytics/export endpoints
  heavyEndpoints: {
    windowMs: 60 * 60 * 1000,   // 1 hour
    max: 20                   // 20 requests per hour
  }
};
```

Offline Capabilities

Offline-First Architecture

Service Worker Implementation

```
// service-worker.js
const CACHE_NAME = 'diagnosisai-v2.0.0';
const STATIC_CACHE = 'diagnosisai-static-v2.0.0';
const DATA_CACHE = 'diagnosisai-data-v2.0.0';

// Static assets to cache
const STATIC_ASSETS = [
  '/',
  '/index.html',
  '/styles.css',
  '/app.js',
  '/manifest.json',
  '/icons/icon-192.png',
  '/icons/icon-512.png'
];

// Install event - cache static assets
self.addEventListener('install', (event) => {
  event.waitUntil(
    caches.open(STATIC_CACHE)
      .then((cache) => cache.addAll(STATIC_ASSETS))
      .then(() => self.skipWaiting())
  );
});

// Activate event - clean up old caches
self.addEventListener('activate', (event) => {
  event.waitUntil(
    caches.keys()
      .then((cacheNames) => {
        return Promise.all(
          cacheNames.map((cacheName) => {
            if (cacheName !== STATIC_CACHE && cacheName !== DATA_CACHE) {
              return caches.delete(cacheName);
            }
          })
        );
      })
      .then(() => self.clients.claim())
  );
});

// Fetch event - network first, then cache
self.addEventListener('fetch', (event) => {
```

```

const { request } = event;
const url = new URL(request.url);

// API requests - network first, cache fallback
if (url.pathname.startsWith('/api/')) {
  event.respondWith(
    fetch(request)
      .then((response) => {
        // Clone and cache successful responses
        if (response.ok) {
          const responseClone = response.clone();
          caches.open(DATA_CACHE).then((cache) => {
            cache.put(request, responseClone);
          });
        }
        return response;
      })
      .catch(() => {
        // Network failed, try cache
        return caches.match(request);
      })
  );
}

// Static assets - cache first
else {
  event.respondWith(
    caches.match(request)
      .then((response) => response || fetch(request))
  );
}
});

// Background sync
self.addEventListener('sync', (event) => {
  if (event.tag === 'sync-pending-data') {
    event.waitUntil(syncPendingData());
  }
});

async function syncPendingData() {
  const db = await openIndexedDB();
  const syncQueue = await db.getAll('syncQueue');

  for (const item of syncQueue) {
    try {

```

```
        await fetch(item.url, {
          method: item.method,
          headers: item.headers,
          body: JSON.stringify(item.data)
        });

        // Remove from queue on success
        await db.delete('syncQueue', item.id);
      } catch (error) {
        console.error('Sync failed for item:', item.id, error);
      }
    }
}
```

Sync Strategy

Queue-Based Synchronization

```

interface SyncQueueItem {
  id: string;
  timestamp: Date;
  operation: 'CREATE' | 'UPDATE' | 'DELETE';
  resource: string;
  resourceId: string;
  data: any;
  attempts: number;
  lastAttempt?: Date;
  status: 'pending' | 'in_progress' | 'failed' | 'completed';
}

class SyncManager {
  private db: OfflineDB;
  private maxRetries = 3;
  private retryDelay = 5000; // 5 seconds

  constructor() {
    this.db = new OfflineDB();
    this.startPeriodicSync();
  }

  async queueForSync(item: Omit<SyncQueueItem, 'id' | 'timestamp' | 'attempts' | 'status'>) {
    const queueItem: SyncQueueItem = {
      id: uuidv4(),
      timestamp: new Date(),
      attempts: 0,
      status: 'pending',
      ...item
    };

    await this.db.put('syncQueue', queueItem);

    // Try to sync immediately if online
    if (navigator.onLine) {
      await this.processSyncQueue();
    }
  }

  async processSyncQueue() {
    const queue = await this.db.getAll('syncQueue');
    const pending = queue.filter(item =>
      item.status === 'pending' && item.attempts < this.maxRetries
    );
  }
}

```

```

for (const item of pending) {
  try {
    // Mark as in progress
    item.status = 'in_progress';
    await this.db.put('syncQueue', item);

    // Execute sync
    await this.syncItem(item);

    // Remove from queue on success
    await this.db.delete('syncQueue', item.id);
  } catch (error) {
    // Update retry info
    item.attempts++;
    item.lastAttempt = new Date();
    item.status = item.attempts >= this.maxRetries ? 'failed' : 'pending';
    await this.db.put('syncQueue', item);

    console.error(`Sync failed for item ${item.id}:`, error);
  }
}

private async syncItem(item: SyncQueueItem) {
  const url = `/api/v1/${item.resource}${item.resourceId ? `/${item.resourceId}` : ''}`;
  const method = item.operation === 'CREATE' ? 'POST' :
    item.operation === 'UPDATE' ? 'PUT' : 'DELETE';

  const response = await fetch(url, {
    method,
    headers: {
      'Content-Type': 'application/json',
      'Authorization': `Bearer ${await this.getAuthToken()}`,
    },
    body: item.operation !== 'DELETE' ? JSON.stringify(item.data) : undefined
  });

  if (!response.ok) {
    throw new Error(`Sync failed: ${response.statusText}`);
  }

  return await response.json();
}

```

```
private startPeriodicSync() {
    // Check for pending syncs every 30 seconds
    setInterval(async () => {
        if (navigator.onLine) {
            await this.processSyncQueue();
        }
    }, 30000);

    // Also sync when coming back online
    window.addEventListener('online', () => {
        this.processSyncQueue();
    });
}

async getSyncStatus() {
    const queue = await this.db.getAll('syncQueue');
    return {
        pending: queue.filter(i => i.status === 'pending').length,
        inProgress: queue.filter(i => i.status === 'in_progress').length,
        failed: queue.filter(i => i.status === 'failed').length
    };
}
}
```

Conflict Resolution

```

interface ConflictResolution {
  strategy: 'last_write_wins' | 'manual' | 'merge';
  resolvedBy?: string;
  resolvedAt?: Date;
}

class ConflictResolver {
  async resolveConflict(
    localVersion: any,
    serverVersion: any,
    strategy: ConflictResolution['strategy'] = 'last_write_wins'
  ) {
    switch (strategy) {
      case 'last_write_wins':
        return this.lastWriteWins(localVersion, serverVersion);

      case 'merge':
        return this.mergeVersions(localVersion, serverVersion);

      case 'manual':
        return this.promptUserResolution(localVersion, serverVersion);

      default:
        throw new Error(`Unknown conflict resolution strategy: ${strategy}`);
    }
  }

  private lastWriteWins(local: any, server: any) {
    const localTime = new Date(local.updatedAt).getTime();
    const serverTime = new Date(server.updatedAt).getTime();

    return serverTime > localTime ? server : local;
  }

  private mergeVersions(local: any, server: any) {
    // Deep merge with preference for non-null server values
    return {
      ...local,
      ...server,
      // Keep whichever is more recent for timestamp fields
      updatedAt: new Date(Math.max(
        new Date(local.updatedAt).getTime(),
        new Date(server.updatedAt).getTime()
      ))
    };
  }
}

```

```
};

}

private async promptUserResolution(local: any, server: any) {
    // Show UI for user to choose
    return new Promise((resolve) => {
        // Implementation would show a modal with both versions
        // and let user choose or merge manually
    });
}
```

Integration Layer

EMR Integration

HL7 FHIR Support

```
import { Client } from 'fhir-kit-client';

class FHIRIntegration {
    private client: any;

    constructor(baseUrl: string) {
        this.client = new Client({ baseUrl });
    }

    // Send patient data to EMR
    async createPatient(patient: Patient) {
        const fhirPatient = this.toFHIRPatient(patient);

        return await this.client.create({
            resourceType: 'Patient',
            body: fhirPatient
        });
    }

    // Send clinical note
    async createDocumentReference(note: ClinicalNote) {
        const documentRef = {
            resourceType: 'DocumentReference',
            status: 'current',
            type: {
                coding: [
                    {
                        system: 'http://loinc.org',
                        code: '11488-4',
                        display: 'Consultation note'
                    }
                ],
                subject: {
                    reference: `Patient/${note.patientId}`
                },
                author: [
                    {
                        reference: `Practitioner/${note.doctorId}`
                    }
                ],
                content: [
                    {
                        attachment: {
                            contentType: 'text/plain',
                            data: Buffer.from(note.fullNote).toString('base64')
                        }
                    }
                ]
            };
    }
}
```

```
        return await this.client.create({
          resourceType: 'DocumentReference',
          body: documentRef
        });
      }

      // Retrieve lab results
      async getLabResults(patientId: string) {
        const response = await this.client.search({
          resourceType: 'Observation',
          searchParams: {
            patient: patientId,
            category: 'laboratory'
          }
        });

        return response.entry.map(entry => this.fromFHIRObservation(entry.resource));
      }

      private toFHIRPatient(patient: Patient) {
        return {
          resourceType: 'Patient',
          identifier: [
            {
              system: 'http://diagnosisai.com/patient-id',
              value: patient.id
            },
            name: [
              {
                family: patient.lastName,
                given: [patient.firstName]
              },
              gender: patient.gender,
              birthDate: patient.dateOfBirth.toISOString().split('T')[0],
              telecom: [
                {
                  system: 'phone',
                  value: patient.contactInfo.phone
                },
                {
                  system: 'email',
                  value: patient.contactInfo.email
                }
              ]
            ];
          };
        }
      }
    }
  }
}
```

```
private fromFHIRObservation(observation: any): LabResult {
  return {
    id: observation.id,
    testName: observation.code.text,
    value: observation.valueQuantity.value,
    unit: observation.valueQuantity.unit,
    referenceRange: {
      low: observation.referenceRange[0].low.value,
      high: observation.referenceRange[0].high.value,
      unit: observation.referenceRange[0].low.unit
    },
    performedAt: new Date(observation.effectiveDateTime),
    // ... additional mapping
  };
}
}
```

Epic Integration (Proprietary)

```
class EpicIntegration {  
  private baseUrl: string;  
  private clientId: string;  
  private clientSecret: string;  
  
  constructor(config: any) {  
    this.baseUrl = config.baseUrl;  
    this.clientId = config.clientId;  
    this.clientSecret = config.clientSecret;  
  }  
  
  async authenticate() {  
    const response = await fetch(`.${this.baseUrl}/oauth2/token`, {  
      method: 'POST',  
      headers: {  
        'Content-Type': 'application/x-www-form-urlencoded'  
      },  
      body: new URLSearchParams({  
        grant_type: 'client_credentials',  
        client_id: this.clientId,  
        client_secret: this.clientSecret  
      })  
    });  
  
    const data = await response.json();  
    return data.access_token;  
  }  
  
  async sendClinicalNote(note: ClinicalNote) {  
    const token = await this.authenticate();  
  
    const response = await fetch(`.${this.baseUrl}/api/FHIR/R4/DocumentReference`, {  
      method: 'POST',  
      headers: {  
        'Authorization': `Bearer ${token}`,  
        'Content-Type': 'application/fhir+json'  
      },  
      body: JSON.stringify({  
        resourceType: 'DocumentReference',  
        // ... Epic-specific formatting  
      })  
    });  
  
    return await response.json();  
  }  
}
```

```
    }  
}
```

Lab System Integration

HL7 v2.x Messaging

```
import * as hl7 from 'simple-hl7';

class HL7Integration {
  async parseLabResult(hl7Message: string): Promise<LabResult> {
    const parser = new hl7.Parser();
    const message = parser.parse(hl7Message);

    // Extract OBX segments (observation/result)
    const obxSegments = message.segments.filter(s => s.name === 'OBX');

    return obxSegments.map(segment => ({
      testName: segment.fields[3],
      value: segment.fields[5],
      unit: segment.fields[6],
      referenceRange: this.parseReferenceRange(segment.fields[7]),
      abnormalFlag: segment.fields[8],
      observationDateTime: new Date(segment.fields[14])
    }));
  }

  createLabOrder(order: LabOrder): string {
    const message = new hl7.Message();

    // MSH - Message Header
    message.addSegment([
      'MSH',
      '^~\\&',
      'DiagnosisAI',
      'DiagnosisAI',
      'LabSystem',
      'LabSystem',
      this.formatHL7DateTime(new Date()),
      '',
      'ORM^001',
      this.generateMessageId(),
      'P',
      '2.5'
    ]);

    // PID - Patient Identification
    message.addSegment([
      'PID',
      '1',
      order.patientId,
```

```

        // ... additional patient fields
    ]);

    // ORC - Common Order
    message.addSegment([
        'ORC',
        'NW',
        order.id,
        '',
        '',
        '',
        '',
        this.formatHL7DateTime(order.orderedAt),
        // ... additional order fields
    ]);

    // OBR - Observation Request
    message.addSegment([
        'OBR',
        '1',
        order.id,
        '',
        order.loincCode,
        order.priority,
        // ... additional request fields
    ]);

    return message.toString();
}

private formatHL7DateTime(date: Date): string {
    return date.toISOString()
        .replace(/[:-]/g, '')
        .replace('T', '')
        .split('.')[0];
}

private generateMessageId(): string {
    return `DIA${Date.now()}${Math.random().toString(36).substr(2, 5).toUpperCase()}`;
}
}

```

Performance Optimization

Frontend Optimization

Code Splitting

```
// Lazy load components
const PatientProfile = () => import('./components/PatientProfile.js');
const Analytics = () => import('./components/Analytics.js');
const Settings = () => import('./components/Settings.js');

// Route-based code splitting
const routes = {
  'home': () => import('./screens/HomeScreen.js'),
  'patient': () => import('./screens/PatientScreen.js'),
  'analytics': () => import('./screens/AnalyticsScreen.js'),
  'settings': () => import('./screens/SettingsScreen.js')
};

async function navigate(route) {
  const module = await routes[route]();
  module.render();
}
```

Image Optimization

```
// Responsive images
function generateResponsiveImage(src, sizes) {
  const srcset = sizes.map(size => {
    const url = `${src}?w=${size}&q=80&fm=webp`;
    return `${url} ${size}w`;
  }).join(', ');

  return {
    src: `${src}?w=${sizes[0]}&q=80&fm=webp`,
    srcset,
    sizes: '(max-width: 768px) 100vw, 50vw'
  };
}

// Lazy loading images
const imageObserver = new IntersectionObserver((entries) => {
  entries.forEach(entry => {
    if (entry.isIntersecting) {
      const img = entry.target;
      img.src = img.dataset.src;
      img.classList.remove('lazy');
      imageObserver.unobserve(img);
    }
  });
});

document.querySelectorAll('img.lazy').forEach(img => {
  imageObserver.observe(img);
});
```

Virtual Scrolling

```
// For large patient lists
class VirtualScroller {
  constructor(container, itemHeight, renderItem) {
    this.container = container;
    this.itemHeight = itemHeight;
    this.renderItem = renderItem;
    this.scrollTop = 0;
    this.containerHeight = container.clientHeight;

    this.container.addEventListener('scroll', () => this.handleScroll());
  }

  setItems(items) {
    this.items = items;
    this.totalHeight = items.length * this.itemHeight;
    this.visibleCount = Math.ceil(this.containerHeight / this.itemHeight) + 1;
    this.render();
  }

  handleScroll() {
    this.scrollTop = this.container.scrollTop;
    this.render();
  }

  render() {
    const startIndex = Math.floor(this.scrollTop / this.itemHeight);
    const endIndex = Math.min(startIndex + this.visibleCount, this.items.length);

    const offsetY = startIndex * this.itemHeight;

    const html = `
      <div style="height: ${this.totalHeight}px; position: relative;">
        <div style="transform: translateY(${offsetY}px);">
          ${this.items.slice(startIndex, endIndex)
            .map(item => this.renderItem(item))
            .join('')}
        </div>
      </div>
    `;

    this.container.innerHTML = html;
  }
}
```

Backend Optimization

Database Indexing

```
-- Patient search index
CREATE INDEX idx_patients_name ON patients (first_name, last_name);
CREATE INDEX idx_patients_dob ON patients (date_of_birth);

-- Visit queries
CREATE INDEX idx_visits_patient ON visits (patient_id, timestamp DESC);
CREATE INDEX idx_visits_doctor ON visits (doctor_id, timestamp DESC);
CREATE INDEX idx_visits_status ON visits (status, timestamp DESC);

-- Diagnosis queries
CREATE INDEX idx_diagnoses_visit ON diagnoses (visit_id);
CREATE INDEX idx_diagnoses_icd10 ON diagnoses (icd10_code);

-- Lab results
CREATE INDEX idx_labs_patient ON lab_orders (patient_id, ordered_at DESC);
CREATE INDEX idx_labs_status ON lab_orders (status, priority);

-- Full-text search on symptoms
CREATE INDEX idx_symptoms_fulltext ON visits USING GIN (to_tsvector('english', chief_complaint));

-- Analytics queries
CREATE INDEX idx_visits_date ON visits (DATE(timestamp));
CREATE INDEX idx_diagnoses_date_category ON diagnoses (DATE(created_at), icd10_code);
```

Query Optimization

```
// Use connection pooling
import { Pool } from 'pg';

const pool = new Pool({
  host: process.env.DB_HOST,
  database: process.env.DB_NAME,
  user: process.env.DB_USER,
  password: process.env.DB_PASSWORD,
  max: 20,                      // Maximum pool size
  idleTimeoutMillis: 30000,
  connectionTimeoutMillis: 2000
});

// Prepared statements
const getPatientQuery = {
  name: 'get-patient',
  text: 'SELECT * FROM patients WHERE id = $1',
  values: (id) => [id]
};

// Use transactions for consistency
async function createVisitWithDiagnosis(visitData, diagnosisData) {
  const client = await pool.connect();

  try {
    await client.query('BEGIN');

    const visitResult = await client.query(
      'INSERT INTO visits (patient_id, doctor_id, ...) VALUES ($1, $2, ...) RETURNING *',
      [visitData.patientId, visitData.doctorId, ...]
    );

    const visit = visitResult.rows[0];

    await client.query(
      'INSERT INTO diagnoses (visit_id, name, icd10_code, ...) VALUES ($1, $2, $3, ...)',
      [visit.id, diagnosisData.name, diagnosisData.icd10Code, ...]
    );
  }

  await client.query('COMMIT');
  return visit;
} catch (error) {
  await client.query('ROLLBACK');
  throw error;
}
```

```
    } finally {
        client.release();
    }
}
```

Caching Strategy

```
import Redis from 'ioredis';

const redis = new Redis({
  host: process.env.REDIS_HOST,
  port: process.env.REDIS_PORT,
  password: process.env.REDIS_PASSWORD,
  retryStrategy: (times) => Math.min(times * 50, 2000)
});

class CacheService {
  // Cache patient data
  async cachePatient(patient: Patient, ttl: number = 3600) {
    const key = `patient:${patient.id}`;
    await redis.setex(key, ttl, JSON.stringify(patient));
  }

  async getCachedPatient(patientId: string): Promise<Patient | null> {
    const key = `patient:${patientId}`;
    const cached = await redis.get(key);
    return cached ? JSON.parse(cached) : null;
  }

  // Cache-aside pattern
  async getPatient(patientId: string): Promise<Patient> {
    // Try cache first
    let patient = await this.getCachedPatient(patientId);

    if (!patient) {
      // Cache miss - fetch from database
      patient = await database.findPatientById(patientId);

      if (patient) {
        // Cache for future requests
        await this.cachePatient(patient);
      }
    }

    return patient;
  }

  // Invalidate cache on update
  async invalidatePatient(patientId: string) {
    await redis.del(`patient:${patientId}`);
  }
}
```

```
// Cache API responses
async cacheResponse(key: string, data: any, ttl: number = 300) {
    await redis.setex(key, ttl, JSON.stringify(data));
}

async getCachedResponse(key: string) {
    const cached = await redis.get(key);
    return cached ? JSON.parse(cached) : null;
}
}
```

Background Jobs

```
import Bull from 'bull';

// Create job queues
const analyticsQueue = new Bull('analytics', {
  redis: {
    host: process.env.REDIS_HOST,
    port: process.env.REDIS_PORT
  }
});

const syncQueue = new Bull('sync', {
  redis: {
    host: process.env.REDIS_HOST,
    port: process.env.REDIS_PORT
  }
});

// Process analytics calculation in background
analyticsQueue.process(async (job) => {
  const { doctorId, period, startDate, endDate } = job.data;

  const analytics = await calculateAnalytics(doctorId, period, startDate, endDate);

  // Cache the result
  await cacheService.cacheResponse(
    `analytics:${doctorId}:${period}:${startDate}`,
    analytics,
    3600
  );

  return analytics;
});

// Schedule analytics generation
async function scheduleAnalyticsGeneration(doctorId: string) {
  // Daily analytics
  await analyticsQueue.add(
    { doctorId, period: 'daily', startDate: new Date(), endDate: new Date() },
    { repeat: { cron: '0 0 * * *' } } // Midnight every day
  );

  // Weekly analytics
  await analyticsQueue.add(
    { doctorId, period: 'weekly', startDate: getWeekStart(), endDate: new Date() },
    { repeat: { cron: '0 0 * * 0' } } // Sunday at midnight
  );
}
```

```
{ repeat: { cron: '0 0 * * 0' } } // Sunday midnight
);
}
```

Deployment Strategy

CI/CD Pipeline

```
# .github/workflows/deploy.yml
name: Deploy DiagnosisAI

on:
  push:
    branches: [main]
  pull_request:
    branches: [main]

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - uses: actions/setup-node@v3
        with:
          node-version: '20'
      - run: npm ci
      - run: npm test
      - run: npm run lint

  build:
    needs: test
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - run: docker build -t diagnosisai/api:${{ github.sha }} .
      - run: docker push diagnosisai/api:${{ github.sha }}

  deploy-staging:
    needs: build
    runs-on: ubuntu-latest
    if: github.ref == 'refs/heads/main'
    steps:
      - uses: azure/k8s-set-context@v1
        with:
          kubeconfig: ${{ secrets.KUBE_CONFIG_STAGING }}
      - run: kubectl set image deployment/api api=diagnosisai/api:${{ github.sha }}

  deploy-production:
    needs: deploy-staging
    runs-on: ubuntu-latest
    if: github.ref == 'refs/heads/main'
    environment:
```

```
name: production
url: https://api.diagnosisai.com
steps:
- uses: azure/k8s-set-context@v1
  with:
    kubeconfig: ${{ secrets.KUBE_CONFIG_PROD }}
- run: kubectl set image deployment/api api=diagnosisai/api:${{ github.sha }}
```

Environment Configuration

```
# Staging
staging:
  database:
    host: staging-db.diagnosisai.com
    name: diagnosisai_staging
    replicas: 1
  redis:
    host: staging-redis.diagnosisai.com
resources:
  api:
    replicas: 2
    cpu: "500m"
    memory: "512Mi"

# Production
production:
  database:
    host: prod-db.diagnosisai.com
    name: diagnosisai_prod
    replicas: 3
  redis:
    host: prod-redis.diagnosisai.com
    cluster: true
resources:
  api:
    replicas: 5
    cpu: "1000m"
    memory: "2Gi"
    autoscaling:
      minReplicas: 5
      maxReplicas: 20
      targetCPU: 70
```

Blue-Green Deployment

```
# blue deployment (current)
apiVersion: apps/v1
kind: Deployment
metadata:
  name: api-blue
spec:
  replicas: 5
  selector:
    matchLabels:
      app: api
      version: blue
  template:
    metadata:
      labels:
        app: api
        version: blue
    spec:
      containers:
        - name: api
          image: diagnosisai/api:1.0.0

---
# green deployment (new)
apiVersion: apps/v1
kind: Deployment
metadata:
  name: api-green
spec:
  replicas: 5
  selector:
    matchLabels:
      app: api
      version: green
  template:
    metadata:
      labels:
        app: api
        version: green
    spec:
      containers:
        - name: api
          image: diagnosisai/api:2.0.0

---
```

```
# Service switches between blue/green
apiVersion: v1
kind: Service
metadata:
  name: api
spec:
  selector:
    app: api
    version: blue # Switch to 'green' for new deployment
  ports:
  - port: 80
    targetPort: 3000
```

Scalability Plan

Horizontal Scaling

Auto-scaling Rules

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: api-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: api
  minReplicas: 5
  maxReplicas: 50
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70
    - type: Resource
      resource:
        name: memory
        target:
          type: Utilization
          averageUtilization: 80
    - type: Pods
      pods:
        metric:
          name: http_requests_per_second
        target:
          type: AverageValue
          averageValue: "1000"
```

Database Scaling

Read Replicas

```
# Primary database (read-write)
Primary:
  host: db-primary.diagnosisai.com
  role: master
  connections: 100

# Read replicas (read-only)
Replicas:
  - host: db-replica-1.diagnosisai.com
    role: replica
    lag_threshold: 100ms
  - host: db-replica-2.diagnosisai.com
    role: replica
    lag_threshold: 100ms
  - host: db-replica-3.diagnosisai.com
    role: replica
    lag_threshold: 100ms
```

Connection Router

```
class DatabaseRouter {
  private primary: Pool;
  private replicas: Pool[];
  private currentReplica: number = 0;

  async query(sql: string, params: any[]) {
    // Write operations go to primary
    if (this.isWriteOperation(sql)) {
      return await this.primary.query(sql, params);
    }

    // Read operations distributed across replicas
    const replica = this.getNextReplica();
    return await replica.query(sql, params);
  }

  private isWriteOperation(sql: string): boolean {
    const writeKeywords = ['INSERT', 'UPDATE', 'DELETE', 'CREATE', 'ALTER', 'DROP'];
    const upperSQL = sql.trim().toUpperCase();
    return writeKeywords.some(keyword => upperSQL.startsWith(keyword));
  }

  private getNextReplica(): Pool {
    // Round-robin load balancing
    const replica = this.replicas[this.currentReplica];
    this.currentReplica = (this.currentReplica + 1) % this.replicas.length;
    return replica;
  }
}
```

CDN & Edge Caching

```
CloudFront Distribution:  
Origins:  
  - API: api.diagnosisai.com  
  - Static Assets: s3://diagnosisai-assets  
  
Behaviors:  
  - PathPattern: /api/*  
    CachingDisabled: true  
    AllowedMethods: [GET, POST, PUT, DELETE, OPTIONS, HEAD, PATCH]  
  
  - PathPattern: /assets/*  
    CachingEnabled: true  
    TTL: 31536000 # 1 year  
    Compression: true  
  
Edge Locations: All  
  
Cache Policies:  
  static-assets:  
    MinTTL: 86400  
    MaxTTL: 31536000  
    DefaultTTL: 86400  
  
  api-responses:  
    MinTTL: 0  
    MaxTTL: 300  
    DefaultTTL: 0
```

Load Testing Results

Target: 10,000 concurrent users

Results:

Average Response Time: 145ms
95th Percentile: 280ms
99th Percentile: 450ms
Error Rate: 0.02%

Peak Performance:

Requests/Second: 50,000
Database Connections: 450
CPU Usage: 65%
Memory Usage: 72%

Bottlenecks Identified:

- AI API calls (external dependency)
- Complex analytics queries
- Full-text search on large datasets

Optimizations Applied:

- Caching AI responses for similar queries
- Pre-computed analytics snapshots
- Elasticsearch for full-text search

Future Enhancements

Roadmap

Q1 2026

- Multi-language support (Spanish, Hindi, Mandarin)
- Voice-to-text symptom input (whisper API)
- Telemedicine integration (video consultations)

Q2 2026

- Medical image analysis (X-rays, CT scans)
- Wearable device integration (Apple Watch, Fitbit)
- Predictive analytics (readmission risk, disease progression)

Q3 2026

- Natural language querying of medical literature
- Collaborative diagnosis (multi-doctor consultations)
- Patient portal (view their own records, test results)

Q4 2026

- Global deployment (Europe, Asia, South America)
 - Insurance claim automation
 - Clinical trial matching
-

End of Technical Architecture Document